




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Introduction to the Sarus container engine

PRACE HPC Workshop on Containers and Unikernels

Lucas Benedicic, ETH Zurich / CSCS

Alberto Madonna, ETH Zurich / CSCS

July 7, 2021

# Table of Contents: morning

1. 9:00 - 9:20 Introduction to Sarus from a user's perspective  
9:20 - 9:30 Q&A
2. 9:30 - 9:45 Installing Sarus on your system  
9:45 - 10:15 Hands on: installing Sarus on a Debian 10 VM

## 10:15 - 10:45 Break

3. 10:45 - 11:15 MPI examples with OSU benchmarks  
11:15 - 12:00 Hands on:
  - Sarus basic commands
  - MPI exercises

## 12:00 - 13:00 Lunch break

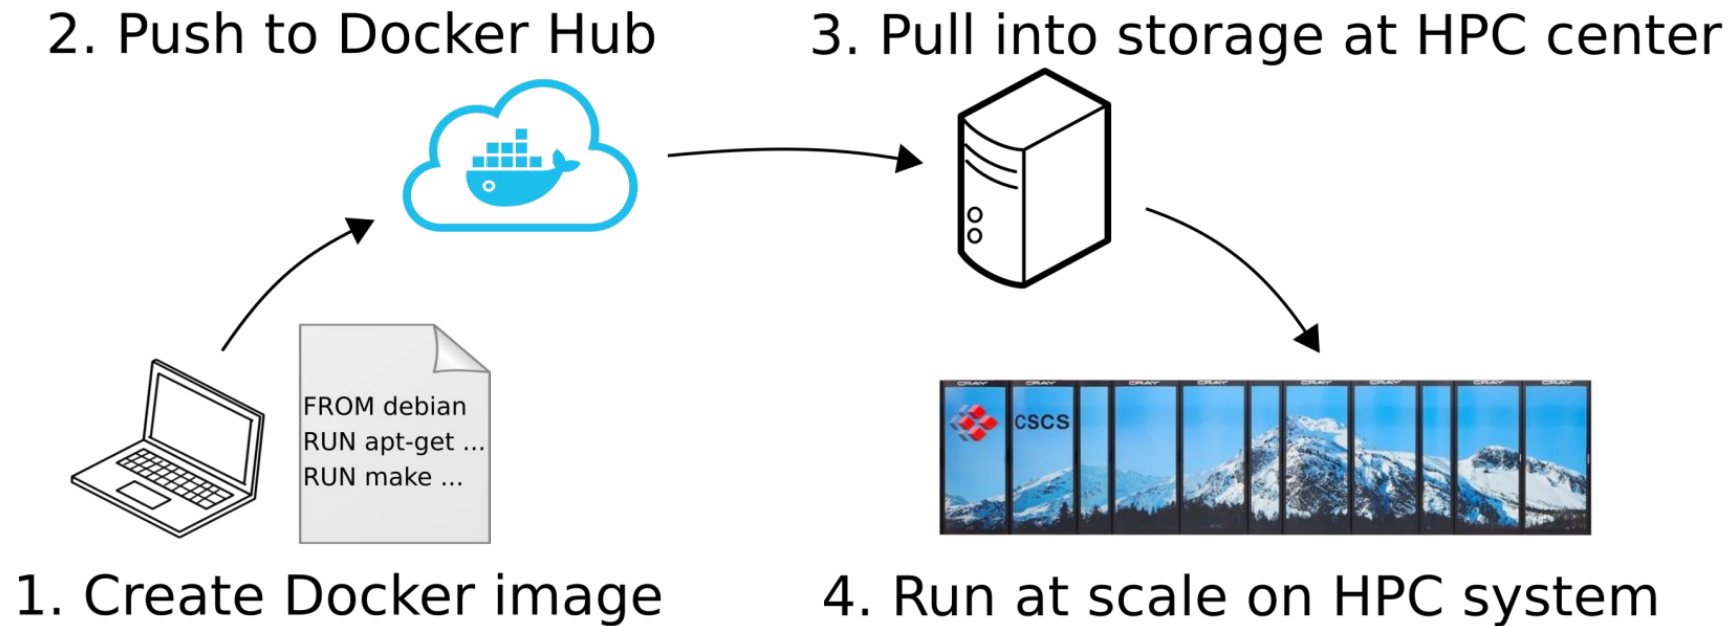
# Sarus container engine



- Combines container portability with native HPC performance
- Integrates with HPC infrastructure and software
- Customizes containers at runtime with standard plugins
- Pulls regular Docker images
- Provides a Docker-like command line interface



# Typical user workflow at CSCS



# Highlights of Sarus from a user perspective

- Consistent experience
  - With Docker: closely resembling CLI
  - With host environment: env variables, uid/gid, file permissions
- Pull images from Docker registries (e.g. Docker Hub, NVIDIA NGC)
- Import images from local tar archives (no cloud upload required)
- Integration with the workload manager (Slurm)
- Native performance from GPUs and high-speed interconnects
- Access to parallel filesystems inside containers

# Sarus CLI

## ■ Sarus

```
# pull image
$ sarus pull [options] <image>[<:tag>]

# load image
$ sarus load [options] <file> <image>

# show list of images
$ sarus images

# remove image
$ sarus rmi <image>[<:tag>]

# run container
$ sarus run [options] <image>[<:tag>]
<command> <args>
```

## ■ Docker

```
# pull image
$ docker pull [options] <image>[<:tag>]

# load image
$ docker load [options] -i <file>

# show list of images
$ docker images [options] [repo[<:tag>]]

# remove image
$ docker rmi [options] <image> [image...]

# run container
$ docker run [options] <image>[<:tag>]
<command> <args>
```

## Further reading

- Sarus on the CSCS User Portal:  
<https://user.cscs.ch/tools/containers/sarus/>
- Code on GitHub:  
<https://github.com/eth-cscs/sarus>
- User documentation on Read the Docs:  
<https://sarus.readthedocs.io/en/stable/user/index.html>
- Benedicic, L., Cruz, F.A., Madonna, A. and Mariotti, K., 2019, June. Sarus: Highly Scalable Docker Containers for HPC Systems. In *International Conference on High Performance Computing* (pp. 46-60). Springer, Cham.  
[https://doi.org/10.1007/978-3-030-34356-9\\_5](https://doi.org/10.1007/978-3-030-34356-9_5)

# Installation tutorial

---



# Hands on!

---

**Break time! Back at 10:45**

---



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# MPI containers on Piz Daint

---



# MPI containers on Piz Daint

- Generic images can run unmodified by instructing Slurm to use the PMI-2 interface:

```
srun --mpi=pmi2 sarus run <image> <args>
```

- This way, containers will use the MPI libraries from the image and run at sub-optimal performance
- Images using MPICH and derivatives: work out of the box
- Images using OpenMPI: OpenMPI must be built with PMI-2 support
  - Configure example on Ubuntu 18.04:

```
./configure --prefix=/usr --with-pmi=/usr/include/slurm-wlm --with-pmi-libdir=/usr/lib/x86_64-linux-gnu \  
CFLAGS=-I/usr/include/slurm-wlm
```

# MPI containers on Piz Daint

- Images using MPICH-based implementations can take advantage of ABI compatibility (<https://www.mpich.org/abi/>)
- Sarus can replace the image MPI with host libraries at runtime, achieving the full performance of the Cray Aries interconnect:

```
srun sarus run --mpi <image> <args>
```

- Recommended libraries for compatibility with Piz Daint:

*MPICH 3.1.4*

*MVAPICH2 2.2*

*Intel MPI Library 2017 Update 1*

# Hands on!

---

**Lunch break! Back at 13:00**

---

## Table of Contents: afternoon

4. 13:00 - 13:30 GPU examples with CUDA SDK  
13:30 - 14:30 Hands on: CUDA samples or user-provided GPU applications

### 14:30 - 15:00 Break

5. 15:00 - 15:30 Real-world application with data I/O: GROMACS  
15:30 - 16:30 Hands on with GROMACS or other applications

6. 16:30 - 17:00 Q&A with the trainers

### 17:00 End of day



# GPU containers on Piz Daint

---

# GPU containers on Piz Daint

- When running on Piz Daint's GPU nodes, GPU devices are automatically added to containers
- Fastest way to get CUDA in a Dockerfile: use NVIDIA official images!  
<https://hub.docker.com/r/nvidia/cuda>

```
FROM nvidia/cuda:11.3.0-devel-ubuntu20.04
```

- NVIDIA images are provided for Ubuntu, Red Hat UBI and CentOS
  - Other distributions can still install the CUDA Toolkit through package manager or runfile
- The NVIDIA driver should NOT be installed in the image (it's bound to the hardware!)

# Hands on!

---

**Break time! Back at 15:00**


---



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**



**EURO**

# Backup slides

---

# Docker and HPC: not a good fit

- Security model assumes root privileges
- No integration with workload managers
- Missing support for diskless nodes
- Very limited support for kernel bypassing devices (e.g. accelerators and NICs)
- No adequate parallel storage driver